

# ShreyaPatil\_MLCaseStudy

July 6, 2025

## 1 Heart Disease Prediction Using Machine Learning

**Student Name :** Shreya Dipak Patil **Course :** Machine Learning using Python **Project Type :** Capstone Project **Objective :** To build a machine learning model that predicts the presence of heart disease in a patient based on clinical features.

### 1.0.1 Step 1: Import essential libraries

```
[34]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

# ML model tools
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import accuracy_score, confusion_matrix, \
    classification_report, roc_curve, roc_auc_score
from sklearn.metrics import ConfusionMatrixDisplay

# ML models
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.svm import SVC
```

### 1.0.2 Step 2: Load the dataset

```
[33]: df = pd.read_csv("heart.csv")
```

```
[4]: # show first 5 rows
df.head()
```

```
[4]:   Age  Sex  ChestPainType  RestingBP  Cholesterol  FastingBS  RestingECG  MaxHR  \
0   40   M             ATA         140          289           0       Normal    172
1   49   F             NAP         160          180           0       Normal    156
2   37   M             ATA         130          283           0           ST     98
```

3	48	F	ASY	138	214	0	Normal	108
4	54	M	NAP	150	195	0	Normal	122

	ExerciseAngina	Oldpeak	ST_Slope	HeartDisease
0	N	0.0	Up	0
1	N	1.0	Flat	1
2	N	0.0	Up	0
3	Y	1.5	Flat	1
4	N	0.0	Up	0

```
[5]: # Shape of the dataset
print("Dataset shape : ",df.shape)
```

Dataset shape : (918, 12)

```
[6]: # Column info and data types
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 918 entries, 0 to 917
Data columns (total 12 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Age                   918 non-null   int64
1   Sex                   918 non-null   object
2   ChestPainType         918 non-null   object
3   RestingBP             918 non-null   int64
4   Cholesterol           918 non-null   int64
5   FastingBS             918 non-null   int64
6   RestingECG           918 non-null   object
7   MaxHR                 918 non-null   int64
8   ExerciseAngina        918 non-null   object
9   Oldpeak               918 non-null   float64
10  ST_Slope              918 non-null   object
11  HeartDisease          918 non-null   int64
dtypes: float64(1), int64(6), object(5)
memory usage: 86.2+ KB
```

```
[7]: # Summary Statistics
df.describe()
```

	Age	RestingBP	Cholesterol	FastingBS	MaxHR	\
count	918.000000	918.000000	918.000000	918.000000	918.000000	
mean	53.510893	132.396514	198.799564	0.233115	136.809368	
std	9.432617	18.514154	109.384145	0.423046	25.460334	
min	28.000000	0.000000	0.000000	0.000000	60.000000	
25%	47.000000	120.000000	173.250000	0.000000	120.000000	
50%	54.000000	130.000000	223.000000	0.000000	138.000000	

75%	60.000000	140.000000	267.000000	0.000000	156.000000
max	77.000000	200.000000	603.000000	1.000000	202.000000


	Oldpeak	HeartDisease
count	918.000000	918.000000
mean	0.887364	0.553377
std	1.066570	0.497414
min	-2.600000	0.000000
25%	0.000000	0.000000
50%	0.600000	1.000000
75%	1.500000	1.000000
max	6.200000	1.000000

```
[8]: # Check for null values
df.isnull().sum()
```

```
[8]: Age          0
     Sex          0
     ChestPainType 0
     RestingBP     0
     Cholesterol   0
     FastingBS     0
     RestingECG    0
     MaxHR         0
     ExerciseAngina 0
     Oldpeak       0
     ST_Slope      0
     HeartDisease  0
     dtype: int64
```

### 1.0.3 Step 3 : Exploratory Data Analysis (EDA)


```
[9]: sns.countplot(x='HeartDisease', data=df)
     plt.title("Heart Disease Class Distribution")
     plt.xlabel("Heart Disease (0 = No, 1 = Yes)")
     plt.ylabel("Count")
     plt.show()
```



```
[10]: # Check how many in each class  
df['HeartDisease'].value_counts()
```

```
[10]: HeartDisease  
1    508  
0    410  
Name: count, dtype: int64
```

```
[13]: # Select only numeric columns  
numeric_df = df.select_dtypes(include=['int64', 'float64'])  
  
plt.figure(figsize=(12,8))  
sns.heatmap(numeric_df.corr(), annot=True, cmap='coolwarm')  
plt.title("Feature Correlation Heatmap")  
plt.show()
```



#### 1.0.4 Step 4 : Preprocessing and Feature Engineering

```
[14]: # Encode categorical columns using one-hot encoding
df_encoded = pd.get_dummies(df, drop_first=True)
df_encoded.head()
```

```
[14]:
```

	Age	RestingBP	Cholesterol	FastingBS	MaxHR	Oldpeak	HeartDisease	\
0	40	140	289	0	172	0.0	0	
1	49	160	180	0	156	1.0	1	
2	37	130	283	0	98	0.0	0	
3	48	138	214	0	108	1.5	1	
4	54	150	195	0	122	0.0	0	

	Sex_M	ChestPainType_ATA	ChestPainType_NAP	ChestPainType_TA	\
0	True	True	False	False	
1	False	False	True	False	
2	True	True	False	False	
3	False	False	False	False	
4	True	False	True	False	

	RestingECG_Normal	RestingECG_ST	ExerciseAngina_Y	ST_Slope_Flat	\
0	True	False	False	False	
1	True	False	False	True	
2	False	True	False	False	
3	True	False	True	True	
4	True	False	False	False	

	ST_Slope_Up
0	True
1	False
2	True
3	False
4	True

```
[15]: # Features and Target
X = df_encoded.drop('HeartDisease', axis=1)
y = df_encoded['HeartDisease']

[16]: # Split : 80% train, 20% test
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳random_state=42)

[17]: scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
```

### 1.0.5 Step 5 : Model Building and Evaluation

I will train and evaluate multiple classification models: - Logistic Regression - Decision Tree - Random Forest - Support Vector Machine (SVM) - K-Nearest Neighbors (KNN)

```
[19]: # Logistic Regression

log_model = LogisticRegression()
log_model.fit(X_train, y_train)
y_pred_log = log_model.predict(X_test)

print("Logistic Regression Results : ")
print("Accuracy : ",accuracy_score(y_test, y_pred_log))
print("Classification Report : \n", classification_report(y_test, y_pred_log))
```

```
Logistic Regression Results :
Accuracy : 0.8532608695652174
Classification Report :
```

	precision	recall	f1-score	support
0	0.80	0.87	0.83	77
1	0.90	0.84	0.87	107

accuracy			0.85	184
macro avg	0.85	0.86	0.85	184
weighted avg	0.86	0.85	0.85	184

[21]: *# Decision Tree*

```
dt_model = DecisionTreeClassifier()
dt_model.fit(X_train, y_train)
y_pred_dt = dt_model.predict(X_test)

print("Decision Tree Results : ")
print("Accuracy : ",accuracy_score(y_test, y_pred_dt))
print("Classification Report : \n", classification_report(y_test, y_pred_dt))
```

Decision Tree Results :

Accuracy : 0.8478260869565217

Classification Report :

	precision	recall	f1-score	support
0	0.83	0.81	0.82	77
1	0.86	0.88	0.87	107

accuracy			0.85	184
macro avg	0.84	0.84	0.84	184
weighted avg	0.85	0.85	0.85	184

[22]: *# Random Forest*

```
rf_model = RandomForestClassifier()
rf_model.fit(X_train, y_train)
y_pred_rf = rf_model.predict(X_test)

print("Random Forest Results : ")
print("Accuracy : ",accuracy_score(y_test, y_pred_rf))
print("Classification Report : \n", classification_report(y_test, y_pred_rf))
```

Random Forest Results :

Accuracy : 0.8641304347826086

Classification Report :

	precision	recall	f1-score	support
0	0.83	0.84	0.84	77
1	0.89	0.88	0.88	107

accuracy			0.86	184
macro avg	0.86	0.86	0.86	184

weighted avg	0.86	0.86	0.86	184
--------------	------	------	------	-----

[23]: *# Support Vector Machine (SVM)*

```
svm_model = SVC()
svm_model.fit(X_train, y_train)
y_pred_svm = svm_model.predict(X_test)

print("SVM Results : ")
print("Accuracy : ",accuracy_score(y_test, y_pred_svm))
print("Classification Report : \n", classification_report(y_test, y_pred_svm))
```

SVM Results :

Accuracy : 0.875

Classification Report :

	precision	recall	f1-score	support
0	0.84	0.87	0.85	77
1	0.90	0.88	0.89	107
accuracy			0.88	184
macro avg	0.87	0.87	0.87	184
weighted avg	0.88	0.88	0.88	184

[24]: *# K-Nearest Neighbors (KNN)*

```
knn_model = KNeighborsClassifier()
knn_model.fit(X_train, y_train)
y_pred_knn = knn_model.predict(X_test)

print("KNN Results : ")
print("Accuracy : ",accuracy_score(y_test, y_pred_knn))
print("Classification Report : \n", classification_report(y_test, y_pred_knn))
```

KNN Results :

Accuracy : 0.8532608695652174

Classification Report :

	precision	recall	f1-score	support
0	0.80	0.87	0.83	77
1	0.90	0.84	0.87	107
accuracy			0.85	184
macro avg	0.85	0.86	0.85	184
weighted avg	0.86	0.85	0.85	184



### 1.0.6 Step 6: Result Comparison and Visualization

I compare accuracy of all trained models and visualize the confusion matrix of the best-performing one.

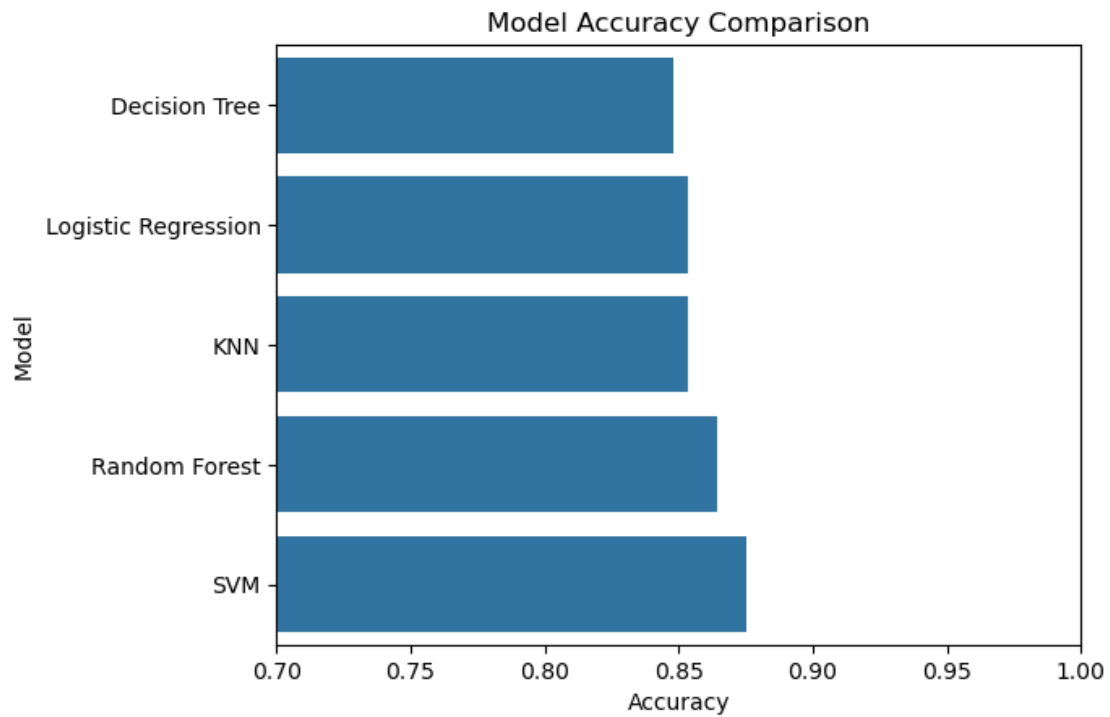
```
[27]: # Compare Accuracy Scores
model_scores = {
    'Logistic Regression' : accuracy_score(y_test, y_pred_log),
    'Decision Tree' : accuracy_score(y_test, y_pred_dt),
    'Random Forest' : accuracy_score(y_test, y_pred_rf),
    'SVM' : accuracy_score(y_test, y_pred_svm),
    'KNN' : accuracy_score(y_test, y_pred_knn)
}

# Show as DataFrame
import pandas as pd
results_df = pd.DataFrame(model_scores.items(), columns=['Model', 'Accuracy'])
results_df.sort_values(by='Accuracy', ascending=False)
```

```
[27]:
```


	Model	Accuracy
3	SVM	0.875000
2	Random Forest	0.864130
0	Logistic Regression	0.853261
4	KNN	0.853261
1	Decision Tree	0.847826

```
[28]: # Plot a Bar Chart
sns.barplot(x='Accuracy', y='Model', data=results_df.sort_values(by='Accuracy',
↪ascending=True))
plt.title('Model Accuracy Comparison')
plt.xlim(0.7, 1)
plt.show()
```



```
[31]: # Confusion Matrix for the best model

cm = confusion_matrix(y_test, y_pred_svm)
disp = ConfusionMatrixDisplay(confusion_matrix=cm)
disp.plot(cmap='Blues')
plt.title("Confusion Matrix - SVM")
plt.show()
```



### Summary of Results

- Among all trained models, **Support Vector Machine (SVM)** gave the highest accuracy.
- **SVM Accuracy: 87.5%**
- The confusion matrix shows that the SVM model correctly classified most positive and negative heart disease cases.
- This makes SVM the best-performing model for this dataset.

### 1.0.7 Step 7: Conclusion

In this project, I built a machine learning model to predict the presence of heart disease based on patient clinical data. We compared five classification algorithms:

- Logistic Regression
- Decision Tree
- Random Forest
- K-Nearest Neighbors

- Support Vector Machine (SVM)

After evaluating all models, **Support Vector Machine (SVM)** performed best with the highest accuracy.

We visualized the confusion matrix of SVM and confirmed its strength in classifying both heart disease and non-disease cases.

#### **1.0.8 Recommendations:**

- More features (like ECG data, lifestyle factors) could improve prediction.
- This model could be deployed in hospitals as a screening tool.